

University of Salford, MSc Data Science

Module: Big Data Tools and Techniques

Date: Trimester 2, 2024-2025

Session: Workshop Week 1

Topic: Linux operating system

Tools: Databricks Community Edition

Instructors: Dr Taha Mansouri, Nathan Topping, and Dr Kaveh Kiani

Objectives:

After completing this workshop, you will be able to:

- Running a Databricks notebook
- Running Some Linux commands
- Working with Local Filesystem

Table of Contents

Part 1: Fire up the Databricks workspace	3
Part 2: Create a Python notebook.....	3
Part 3: LINUX Operating System (OS).....	4
1) What is LINUX	5
2) What is a Filesystem (FS)?.....	6
3) Linux commands	8
Part 4: Running some Linux commands in the Databricks notebook	9
1) How to use Linux commands in a notebook	9
2) Explore a directory.....	10
3) Explore files content.	12
4) Final touch	16
References and Resources:.....	16

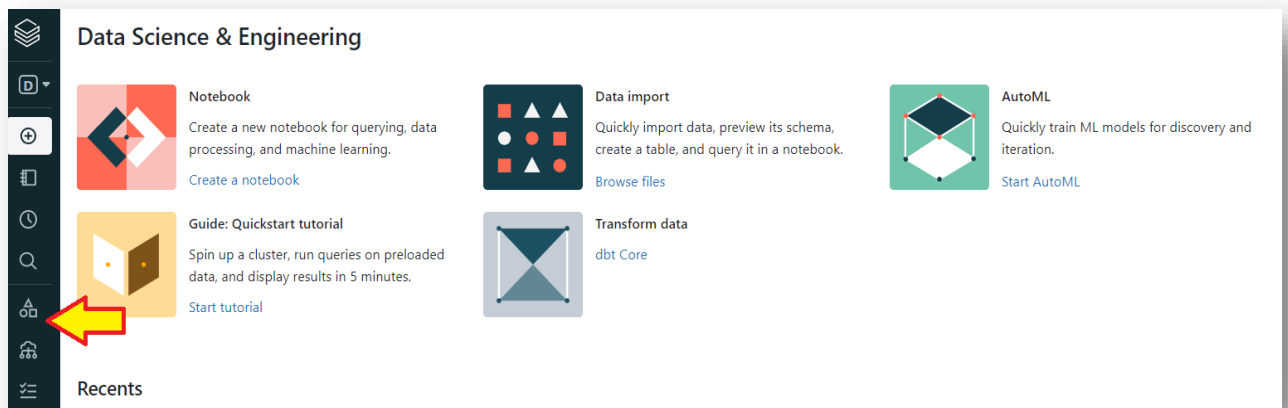
Important Note: In all BDTT workshop notes, if the screenshots appear slightly different from what you see on your screen, there's no need to worry. The Databricks platform is highly dynamic, constantly updating its contents and procedures. Please proceed with your work as usual and we will let you know if any major changes arise.

Part 1: Fire up the Databricks workspace

1. Log in to your Databricks Community Edition account if already you have an account:

<https://community.cloud.databricks.com/login.html>

2. You will need to create a cluster to start your analysis – this gives you access to a machine to use. Click on “Compute” on the main page.



3. Click on “Create Compute” and type in a new name for the cluster in the “compute name” box, any name that you like.
4. Select the most recent runtime (currently, runtime 14.2 (Scala 2.12, Spark 3.5.0) and press create compute

Note that it will take a few minutes to create the cluster. After some time, the green circle next to the cluster name will gain a green tick meaning the cluster has successfully started up.

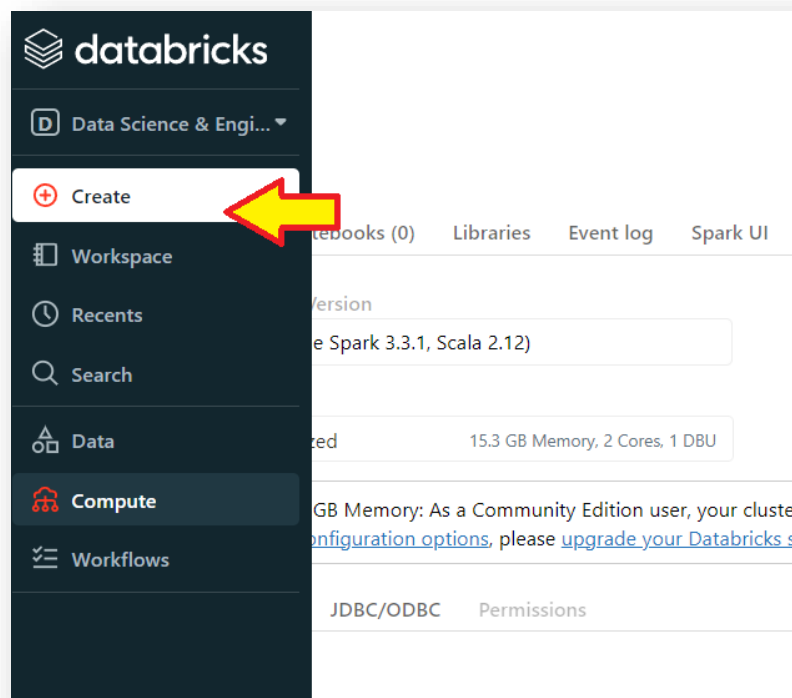
Part 2: Create a Python notebook

Databricks programs are written in Notebooks. A notebook is a collection of runnable cells which contain your commands. Look at

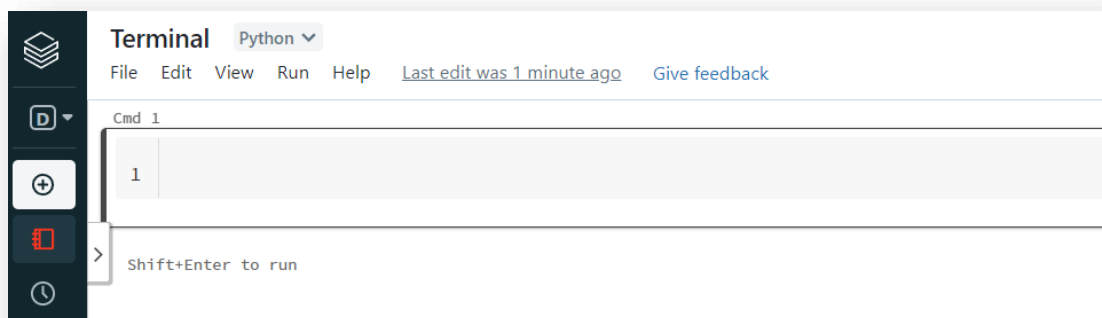
<https://docs.databricks.com/notebooks/notebooks-use.html>

for an overview of how to use a Databricks notebook.

1. Click on the button “Create” and select “Notebook”.



2. Name your notebook, e.g. “Terminal”, leave (or select) the language to be “Python” and the cluster should be your currently created cluster. This will give you access to a notebook which looks something like this:



Part 3: LINUX Operating System (OS)

1) What is LINUX

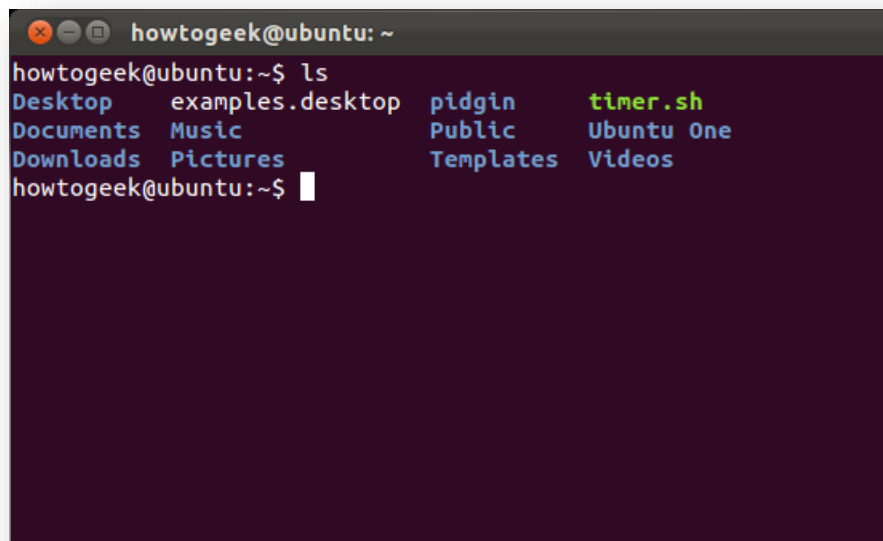
Linux is an open-source **operating system** (OS). So, LINUX is an operating system but unlike the Microsoft Windows, is a free and open-source OS. Data Scientists use Linux because:

- Effective help with quick manipulation and analysis of data
- Data science tools come with a command line interface
- Repeatable systems (more about this later)

When you use a computer, you should work with a OS. Most of us are used to work with Windows or macOS. You know that an **Operating System (OS)** is a software that acts as an interface between computer hardware components and the user. Every computer system must have at least one operating system to run other programs. Applications like Browsers, MS Office, etc. Databricks have assigned you a computer system on the Cloud (The cluster that you created). To communicate with this cluster, you must use Linux OS.


Unlike Windows, in Linux, you need to write commands to interact with the OS. If you need to create a folder, you can't right-click and then click on the New Folder tab, you should write a command. The same applies to Copy and Paste and so on. But where you can write your commands?

The **command line** is an interface that allows you for typing commands directly to a computer's operating system.



```
howtogeek@ubuntu: ~  
howtogeek@ubuntu:~$ ls  
Desktop      examples.desktop  pidgin          timer.sh  
Documents    Music             Public          Ubuntu One  
Downloads    Pictures          Templates       Videos  
howtogeek@ubuntu:~$
```

A Linux basic command is made up of three parts: Name, Options, and Arguments:

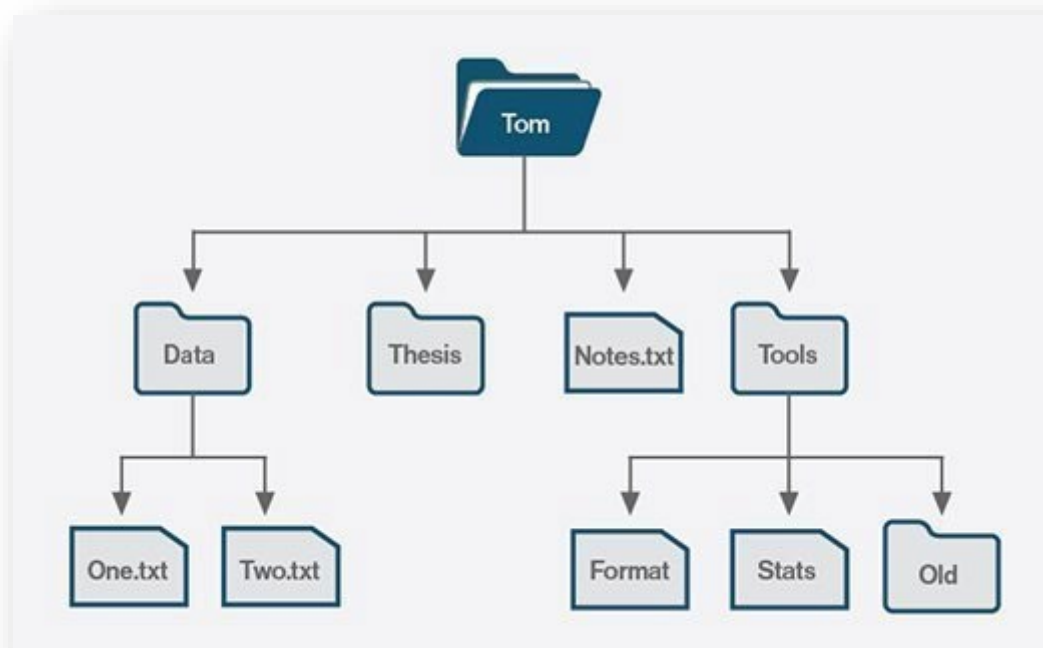


2) What is a Filesystem (FS)?

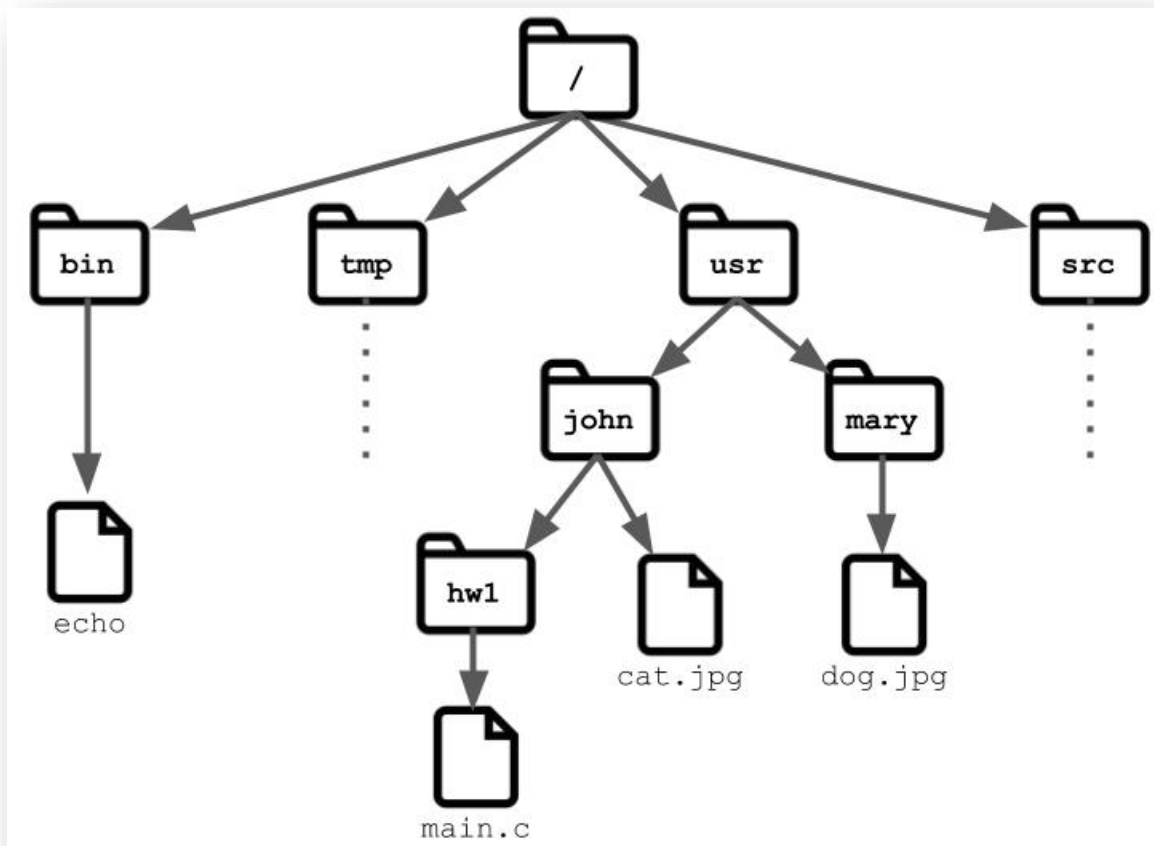
Computers use particular kinds of filesystems (FS) to store and organize data in drives, such as a hard drive or flash drive, or the CDs, DVDs, or Clouds.

A filesystem can be thought of as an index or a database containing the address of the physical location of every piece of data on the device. The data is usually organized in **folders** called **directories**, which can contain other directories and **files**. For example, Windows OS offers three filesystems: NTFS, FAT32 and FAT16.

Imagine Tom uses Windows OS and wants to open/copy/delete the “One.txt” file. Tom must click on the “Tom” folder, then “Data” folder and finally do the intended operation on the “One.txt”.



Also, Linux has a filesystem (FS) and to interact with any content that has been stored in the computer you should use its filesystem logic to do so. Hierarchy of the Linux filesystem is like:



In this example:

1) You can see the “**root**”



2) You can see multiple “**directories**”



...

3) You can see multiple “**files**”



...

To get access to each of these entities you should type the address (**path**) of the entity in the command line.

► **Absolute path**

Start at the root and follow the tree
`/bin/echo`


► **Relative path**

Start at working directory
(ex: current directory is `john`)

`hw1/main.c`

► `..` refers to level above

`../mary/dog.jpg`



3) Linux commands

Main Linux commands are:

(You don't need to learn all this commands now, but we will use some of them during the workshops)

File system exploration

- `ls` → shows contents of a directory
- `cd` → change current working directory
- `pwd` → current working directory
- `du` → estimate file space usage
- `df` → report file system disk space usage
- `man` → interface to on-line reference manuals

Data manipulation

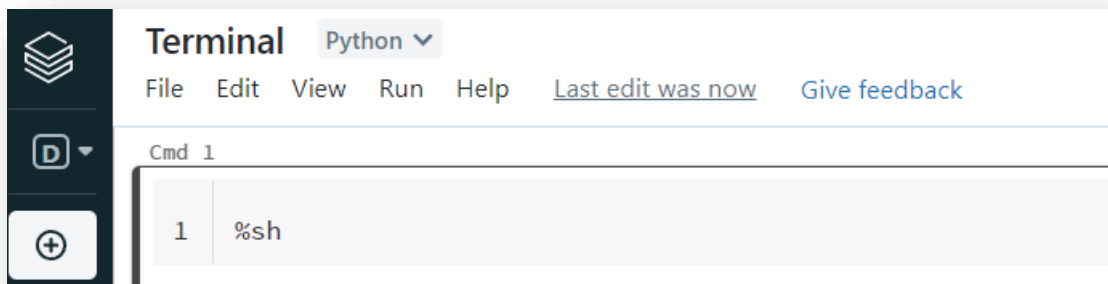
- ▶ `gunzip / gzip` → compress or expand files
- ▶ `curl / wget` → transfer a URL
- ▶ `vim` and `nano` → editors
- ▶ `tr` → translate or delete characters
- ▶ `sort` → sort lines of text files
- ▶ `awk` → pattern scanning and processing language
- ▶ `>`, `>>` → output / append to file
- ▶ `—` pipe → send output to another program

Part 4: Running some Linux commands in the Databricks notebook

1) How to use Linux commands in a notebook

The cells in the notebook expect you to be typing python code, but you can change the type of the cell by using “cell magic”.

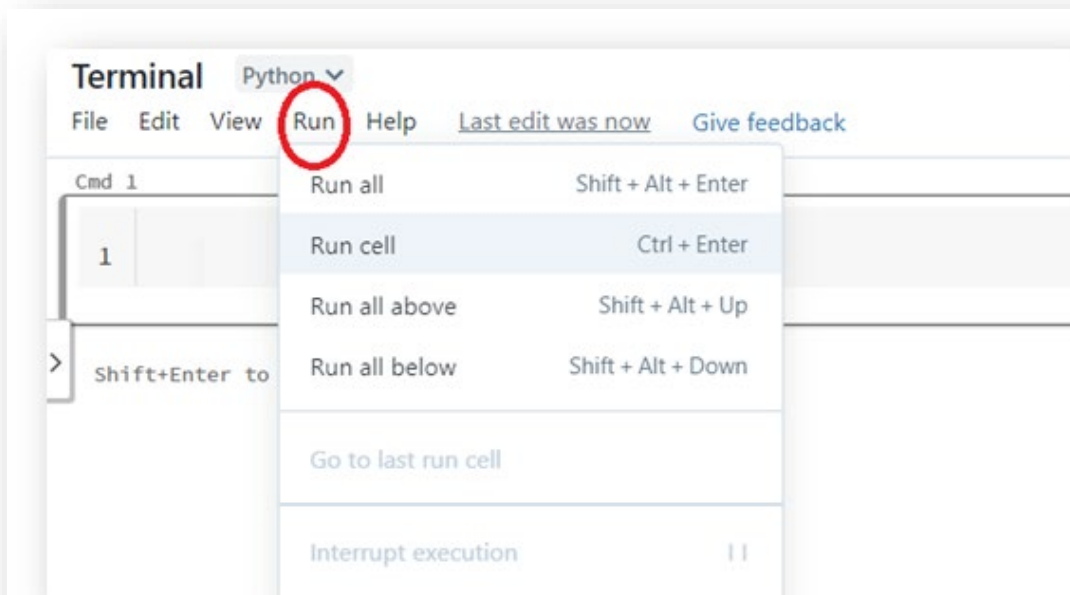
To switch to “shell” and write Linux commands, you can type **%sh** at the start of your cell.



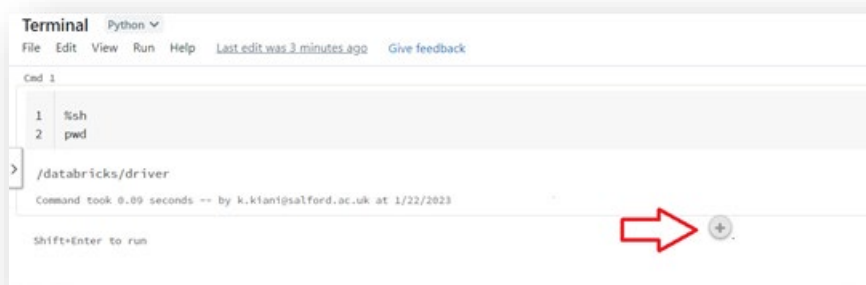
2) Explore a directory

Note 1: To run a cell use Ctrl+Enter or shift+Enter (shift+Enter will open a new cell as well)

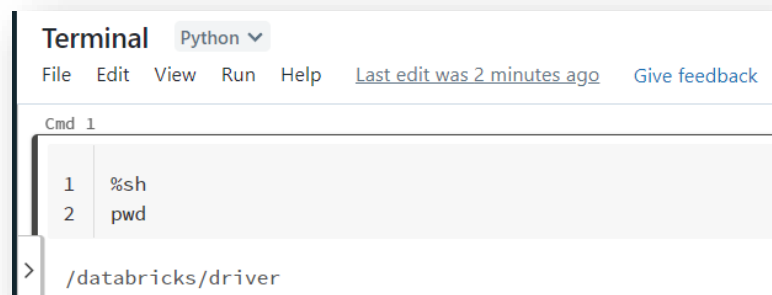
To run a cell or specific cells or all cells, use:



Note 2: To open a new cell press the **+** sign:



1. Check that you are currently in **/databricks/drivers** directory using **pwd** command.



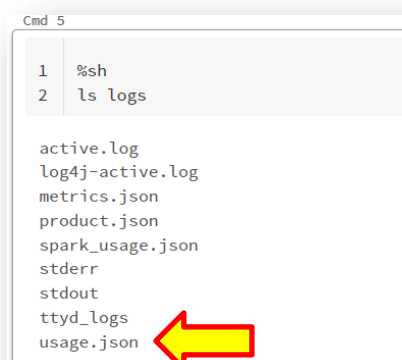
```
Terminal Python ▾  
File Edit View Run Help Last edit was 2 minutes ago Give feedback  
Cmd 1  
1 %sh  
2 pwd  
> /databricks/driver
```

2. As in the cell shown above, you will need to ensure that the top of each cell containing shell commands starts with a **%sh** (this is only required once per cell). Now use the **ls** command to tell you the contents of the directory.



```
Cmd 4  
1 %sh  
2 ls  
  
azure  
conf  
eventlogs  
ganglia  
hadoop_accessed_config.lst  
logs  
metastore_db  
preload_class.lst
```

3. What does **ls -l** in the same place tell you? (You can use Google "ls command Linux" to give you information about the command.)
4. The logs directory looks interesting, look at what's in there using:



```
Cmd 5  
1 %sh  
2 ls logs  
  
active.log  
log4j-active.log  
metrics.json  
product.json  
spark_usage.json  
stderr  
stdout  
ttyd_logs  
usage.json
```

3) Explore files content.

5. What command would you use to view the contents of the “logs/usage.json” file?

```
Cmd 6

1 %sh
2 less logs/usage.json

{"metric":"serviceUptime","tags":{"hostName":"0122-0922
7386ada0000","opType":"TimerTask","opId":"TimerTask-a6d
d":"Thread-12"},"timestamp":1674379381063,"v":3,"quanti
{"metric":"service","tags":{"hostName":"0122-092225-9m2
Name":"","eventType":"start","projectName":"driver"},"t
{"metric":"operationStart","tags":{"hostName":"0122-092
df07386ada0002","opType":"ServiceMain","opId":"ServiceM
ntOpId":"Thread-1"},"timestamp":1674379382231,"v":3,"qu
{"metric":"operationStart","tags":{"hostName":"0122-092
ada0004","branchName":"development","shardName":"","opT
e":"Local","tier":"tier-0","rootOpId":"ServiceMain-a6df
{"metric":"driverMetastoreEvent","tags":{"hostName":"01
ceMain-a6df07386ada0002","branchName":"development","sh
OpId":"ServiceMain-a6df07386ada0002","metastoreType":"u
\":"0.13.0\"}"}
```

6. Extract the values containing the word “metric” from the logs/usage.json file.

```
Cmd 8

1 %sh
2 grep 'metric' logs/usage.json

{"metric":"serviceUptime","tags":{"driverPublicDns":"ec2-18-246-63-110.us-west-2.c
execution_env":"","driverContainerId":"9c73e707330943f2a93a36105511fb9d","instance
LogDeliveryEnabled":"false","sparkEnvVarContainsSingleQuotes":"false","driverNodeT
7","sparkMasterUrlType":"None","userProvidedRemoteVolumeCount":"0","clusterUnityCa
nableJdbcAutoStart":"true","driverInstanceId":"i-0fbcdada2dfa9c51c","orgId":"22801
CatalogCredentialPassthrough":"false","cloudProvider":"AWS","clusterName":"kaveh 1
sterResourceClass":"default","enableSqlAclsOnly":"false","containerType":"LXC","at
6","enableElasticDisk":"false","clusterSizeType":"VM_CONTAINER","dataPlaneRegion":
videdSparkVersion":"11.3.x-scala2.12","clusterEbsVolumeType":"GENERAL_PURPOSE_SSD"
pType":"ssh","awsWorkspaceIMDSV2EnablementStatus":"false","containerZoneId":"us-we
terSku":"STANDARD_SKU","attribute_tag_dust_maintainer":"","ngrokNpipEnabled":"fals
sSingleUserCluster":"false","driverInstancePrivateIp":"10.172.253.26","sparkEnvVar
criptsV2":"0","numPerClusterInitScriptsV2":"0","enableJobsAutostart":"true","cLust
tiveSparkVersion":"11.3.x-scala2.12","projectName":"driver","region":"us-west-2","
s":"false","clusterGeneration":"0","autoTerminationMinutes":"120","clusterId":"012
e","workerEnvironmentId":"default-worker-env","sparkEnvVarContainsNewline":"false"
aemon":"false","clusterOwnerUserId":"8578152565608834","privateLinkEnabled":"false
terSpotBidPricePercent":"100","isDpCpPrivateLinkEnabled":"false","runtimeEngine":
d":"TimerTask-a6df07386ada696f","branchName":"development","userProvidedRemoteVolu
t":"UntimeLogger:driver","clusterMetastoreAccessType":"RDS_DIRECT","attribute_tag_
{"metric":"sparkNumPendingTasks","tags":{"driverPublicDns":"ec2-18-246-63-110.us-w
Command took 2.58 seconds -- by k.kiani@salford.ac.uk at 1/22/2023, 10:33:52 AM on kaveh 1

Cmd 9
```

7. Now we want to extract values of the “metric”.

```
{ "metric": "ServiceUptime", "tags": { "driverPublic  
execution_env": "1", "driverContainerId": "9c73e707  
LogDeliveryEnabled": "false", "sparkEnvVarContain  
7", "sparkMasterUrlType": "None", "userProvidedRe  
nableJdbcAutoStart": "true", "driverInstanceId": "  
CatalogCredentialPassthrough": "false", "cloudPro  
sterResourceClass": "default", "enableSqlAclsOnly  
6", "enableElasticDisk": "false", "clusterSizeType  
videdSparkVersion": "11.3.x-scala2.12", "clusterf  
pType": "ssh", "awsWorkspaceIMDSV2EnablementStatu  
terSku": "STANDARD_SKU", "attribute_tag_dust_mai  
sSingleUserCluster": "false", "driverInstancePriv  
criptsV2": "0", "numPerClusterInitScriptsV2": "0"  
tiveSparkVersion": "11.3.x-scala2.12", "projectNa  
s": "false", "clusterGeneration": "0", "autoTermin  
e", "workerEnvironmentId": "default-worker-env",  
aemon": "false", "clusterOwnerUserId": "8578152565  
terSpotBidPricePercent": "100", "isDpCpPrivateLin  
d": "TimerTask-a6df07386ada696f", "branchName": "c  
t": "UptimeLogger:driver", "clusterMetastoreAcces  
{ "metric": "sparkNumPendingTasks", "tags": { "drive  
Command took 2.58 seconds by M.Kiani@salford.ac.uk a
```

We want to use “cut” command. Run the following help command to learn “cut” functionality.

```
Cmd 9
1 %sh
2 cut --help

Usage: cut OPTION... [FILE]...
Print selected parts of lines from each FILE to standard output.

With no FILE, or when FILE is -, read standard input.

Mandatory arguments to long options are mandatory for short options too.
-b, --bytes=LIST      select only these bytes
-c, --characters=LIST  select only these characters
-d, --delimiter=DELIM use DELIM instead of TAB for field delimiter
-f, --fields=LIST      select only these fields; also print any line
                        that contains no delimiter character, unless
                        the -s option is specified
-n                     (ignored)
--complement           complement the set of selected bytes, characters
                        or fields
-s, --only-delimited   do not print lines not containing delimiters
--output-delimiter=STRING use STRING as the output delimiter
                        the default is to use the input delimiter
-z, --zero-terminated  line delimiter is NUL, not newline
--help                display this help and exit
--version              output version information and exit
```


Here is the command:

```
Cmd 7
1 %sh
2 grep 'metric' logs/usage.json | cut -d '"' -f 4

serviceUptime
sparkNumPendingTasks
sparkNumTaskSlotsTotal
sparkNumActiveTasks
sparkNumExecutorsTotal
sparkNumExecutorsActive
sparkNumExecutorsActiveOrHoldingShuffles
sparkNumPendingTasks
sparkNumTaskSlotsTotal
sparkNumActiveTasks
sparkNumExecutorsTotal
sparkNumExecutorsActive
sparkNumExecutorsActiveOrHoldingShuffles
serviceUptime
```

What is the usage of the vertical bar “|” in between of the code? “|” is the **pipeline** and like a pipe passes result of the first part to the second part. The vertical bar connects the commands together, making it possible to create a chain of related but separate processes.

First part extracts the values containing the word “**metric**” from the logs/usage.json file and then second part receive outputs of the first part and extracts intended part of the inputs.



In the second command (cut -d '"' -f 4) we ask Linux to extract (or cut) based on “” delimiter and give use the forth part of the each object. Let’s take a close look to one of the objects that has been started with the word “metric”.

Start of the first object that contains “metric” word



```
{
  "metric": "serviceUptime",
  "tags": {
    "driverPublicDns": "ec2-18-246-63-110.us-west-2.compute.amazonaws.com",
    "opType": "TimerTask",
    "clusterScalingType": "fixed_size",
    "attribute_tag_dust_execution_env": "",
    "driverContainerId": "9c73e707330943f2a93a36105511fb9d",
    "instanceProfileUsed": "false",
    "clusterState": "Pending",
    "ignoreTerminationEventInAlerting": "false",
    "clusterLogDeliveryEnabled": "false",
    "sparkEnvVarContainsSingleQuotes": "false",
    "driverNodeType": "dev-tier-node",
    "shardName": "",
    "databricksCommit": "b53583c28bddd14ab6e3579938248b013ff01807",
    "sparkMasterUrlType": "None",
    "userProvidedRemoteVolumeCount": "0",
    "clusterUnityCatalogMode": "CUSTOM",
    "hostName": "0122-092225-9m2v72np-10-172-244-26",
    "clusterEbsVolumeSize": "0",
    "enableJdbcAutoStart": "true",
    "driverInstanceId": "i-0fbcdada2dfa9c51c",
    "orgId": "2280122027430776",
    "sparkVersion": "11.3.x-scala2.12",
    "clusterOwnerOrgId": "2280122027430776",
    "enableGlueCatalogCredentialPassthrough": "false",
    "cloudProvider": "AWS",
    "clusterName": "kaveh_1",
    "clusterPythonVersion": "3",
    "clusterAllTags": "[{key: 'Name', value: 'ce-worker'}]",
    "clusterResourceClass": "default",
    "enableSqlAclsOnly": "false",
    "containerType": "LXC",
    "attribute_tag_dust_suite": "",
    "userId": "8578152565608834",
    "driverContainerPrivateIp": "10.172.244.26",
    "enableElasticDisk": "false",
    "clusterSizeType": "VM_CONTAINER",
    "dataPlaneRegion": "us-west-2",
    "instanceWorkerEnvNetworkType": "default",
    "sparkEnvVarContainsEscape": "false",
    "userProvidedSparkVersion": "11.3.x-scala2.12",
    "clusterEbsVolumeType": "GENERAL_PURPOSE_SSD",
    "instanceWorkerEnvId": "default-worker-env",
    "isServicePrincipalCluster": "false",
    "instanceBootstrapType": "ssh",
    "awsWorkspaceIMDSV2EnablementStatus": "false",
    "containerZoneId": "us-west-2c",
    "clusterPinned": "false",
    "clusterNodeType": "dev-tier-node",
    "clusterNumCustomTags": "0",
    "clusterSku": "STANDARD_SKU",
    "attribute_tag_dust_maintainer": "",
    "ngrokNpipEnabled": "false",
    "clusterFirstOnDemand": "0",
    "enableCredentialPassthrough": "false",
    "clusterLogDestination": "",
    "isSingleUserCluster": "false",
    "driverInstancePrivateIp": "10.172.253.26",
    "sparkEnvVarContainsDollarSign": "false",
    "isIMv2Enabled": "false",
    "clusterTargetWorkers": "0",
    "numPerGlobalInitScriptsV2": "0",
    "numPerClusterInitScriptsV2": "0",
    "enableJobsAutostart": "true",
    "clusterStateMessage": "Starting Spark",
    "apacheCommit": "990bee9c58ea9abd8c4f04f20c78c6d5b720406a",
    "effectiveSparkVersion": "11.3.x-scala2.12",
    "projectName": "driver",
    "region": "us-west-2",
    "userProvidedRemoteVolumeSizeGb": "0",
    "attribute_tag_service": "",
    "sparkEnvVarContainsDoubleQuotes": "false",
    "clusterGeneration": "0",
    "autoTerminationMinutes": "120",
    "clusterId": "0122-092225-9m2v72np",
    "clusterAvailability": "ON_DEMAND",
    "parentOpId": "Thread-12",
    "hailEnabled": "false",
    "workerEnvironmentId": "default-worker-env",
    "sparkEnvVarContainsNewline": "false",
    "tier": "tier-0",
    "enableDfAcls": "false",
    "rootOpId": "TimerTask-a6df07386ada696f",
    "clusterNoDriverDaemon": "false",
    "clusterOwnerUserId": "8578152565608834",
    "privateLinkEnabled": "false",
    "clusterEbsVolumeCount": "0",
    "clusterCreator": "Webapp",
    "enableLocalDiskEncryption": "false",
    "clusterSpotBidPricePercent": "100",
    "isDpCpPrivateLinkEnabled": "false",
    "runtimeEngine": "STANDARD",
    "buildHash": "c3dcb270ce6f862df4f4a42d1a751c1a4ba96e4",
    "clusterNumSshKeys": "0",
    "opId": "TimerTask-a6df07386ada696f",
    "branchName": "development",
    "userProvidedRemoteVolumeType": "ebs_volume_type: GENERAL_PURPOSE_SSD/n",
    "sparkEnvVarContainsBacktick": "false",
    "opTarget": "UptimeLogger:driver",
    "clusterMetastoreAccessType": "RDS_DIRECT",
    "attribute_tag_budget": "",
    "clusterWorkers": "0",
    "timestamp": "1674381600172",
    "v": "3",
    "quantity": "30000.0",
    "blob": null,
    "metric": "sparkNumPendingTasks",
    "tags": {
      "driverPublicDns": "ec2-18-246-63-110.us-west-2.compute.amazonaws.com",
      "opType": "ServiceMain",
      "clusterScalingType": "fixed_size",
      "attribute_tag_dust_execution_env": ""
    }
  }
}
```



End of the first object that contains “metric” word

`{"metric": "serviceUptime", "tags":`



We want to extract this string and we have defined “ as delimiter

(cut -d "" -f 4) command says to extract the 4th part of the object where “ is delimiter (separator). Let’s count together:

`{"metric": "serviceUptime", "tags":`



1st



2nd



3rd



4th

Challenge 1:

Now try to extract values for "driverPublicDns:

```
"driverPublicDns": "ec2-18-246-63-110.us-west-2.compute.amazonaws.com"
```

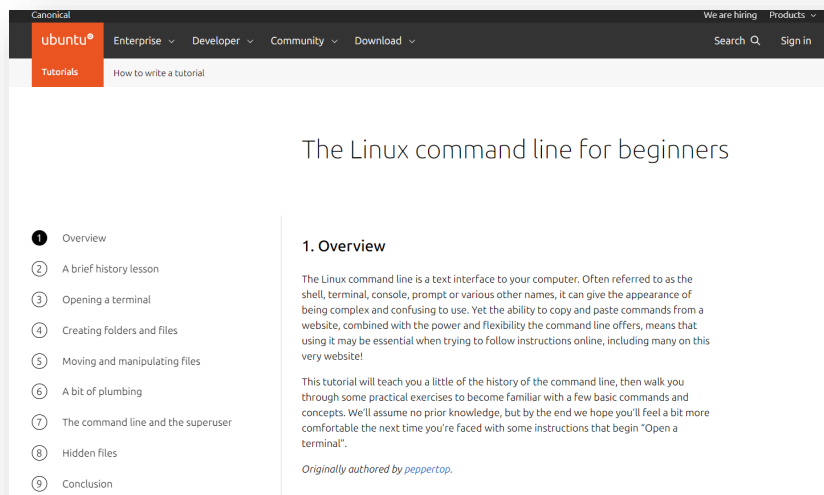
Challenge 2:

Can you think of a way to print each of the values only once?

4) Final touch

Ready to level up? Head over to this link and explore the website. It's a great way to test your skills, refresh your memory and learn more as much as you can and no need to learn every single detail of it for this workshop.

<https://ubuntu.com/tutorials/command-line-for-beginners#1-overview>



References and Resources:

<https://docs.databricks.com/files/index.html>

<https://docs.databricks.com/notebooks/notebooks-use.html>

<https://ubuntu.com/tutorials/command-line-for-beginners#1-overview>