



University of
Salford
MANCHESTER



SCHOOL OF
**SCIENCE, ENGINEERING
& ENVIRONMENT**

Big Data Tools and Techniques

Week 6

Real-Time Magic

Stream Processing & Structured Streaming in
PySpark

2025

Expectations


1. Choose a quiet place to attend the class and please concentrate during the lecture.
2. Put your questions in Padlet and I will review them in the due time (Padlet link is in BB, week 6, Lecture folder for Q&A week6).
3. You can find a handout on BB.
4. We will have 5 mins break after the first hour of the lecture (please remind me).
5. Jisc code will be shared during the break time.

Learning Outcomes

1. To recognize the stream processing
2. To describe differences between batch and stream processing
3. To apply Apache Spark Structured Streaming in different use cases.

The background of the slide is a dense, chaotic web of thin, tangled lines. The lines are primarily shades of pink and red, with some black lines interspersed. The lines are of varying lengths and curves, creating a complex, almost abstract pattern that fills the entire frame. The overall effect is one of intense energy and complexity.


Recap



Different Types of Analysis

Types of Analysis	Questions	Value	Complexity	Example
Descriptive	What happened?	Summarizes past performance.	Low	Monthly sales report.
Diagnostic	Why did it happen?	Identifies causes of past events.	Moderate	Analysing drop in sales for a region.
Predictive	What will happen?	Forecasts future outcomes.	High	Predicting product demand for the next quarter.
Prescriptive	What should we do about it?	Recommends actions to achieve goals.	Very high	Optimizing inventory levels for the predicted demand to avoid stockouts.

Characteristics of Big Data (The 5 Vs)




Data Analytics



Static Batch Data

Allows comprehensive processing at scheduled intervals



Input data



Historical datasets




Modelling



Insight


Question?

Does this approach work for every applications?




Activity



Join at menti.com | use code 6595 8196

 University of Salford

In one word, what does real-time data processing mean to you?



transpiration
bold
inspiration
creative
fast
focus
leader




 TM

Menti

BDTT-S-6-1



Choose a slide to present




Stream Processing




As a barista, you don't pile up the orders.

Stream processing


Unveiling the Core of Stream Processing Pipelines



Real-Time Data Processing Funnel

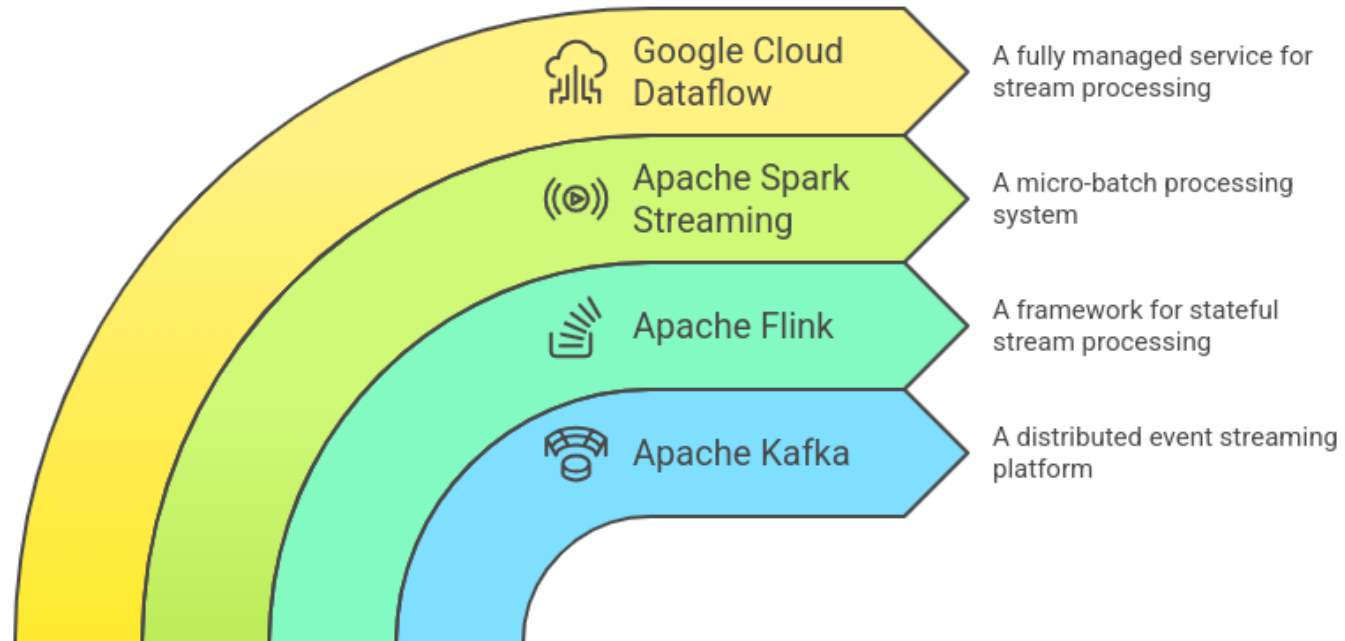


Data Sources




Engines

Stream Processing Engine Overview



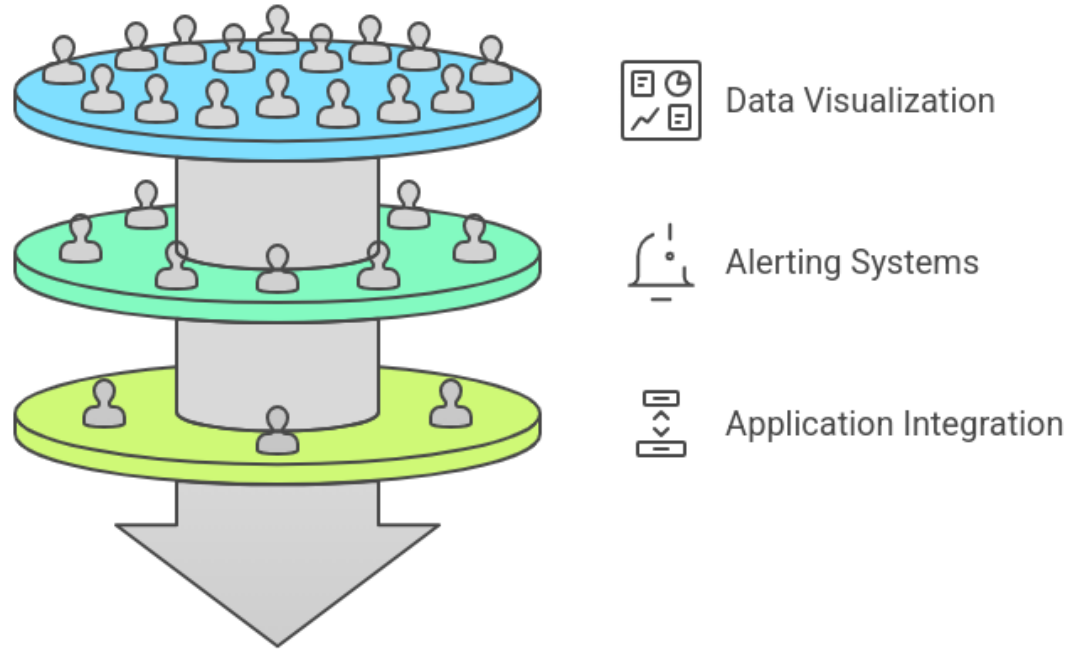
Data Storage

How should processed data be stored?



Data Sinks

Streamlining Data to Actionable Insights



Monitoring and Management

Components of Stream Processing Pipelines



Key Features

- Real-time Processing
- Low Latency
- Data Streams
- Stateful Processing
- Event-Driven
- Windowing

Key Challenges

- Complexity of Handling Late Data
- State Management
- Resource & Performance Tuning
- Increased System Complexity
- Data Quality

News

- [Wes Streeting unveils plans for 'patient passports' to hold all medical records](#)

News

Opinion

Sport

Culture

Lifestyle

UK

The Guardian

UK US politics World Climate crisis Middle East Ukraine Football Newsletters Business Environment UK politics Society Science Tech Global development Obituaries

NHS

This article is more than 3 months old

Wes Streeting unveils plans for 'patient passports' to hold all medical records

Health secretary launches consultation on government's move to transform NHS in England from 'analogue to digital'

Pippa Crerar and Denis Cammell

Advertisement



CAPTAIN AMERICA
BRAVE NEW WORLD

AT ODEON NOW

EXPERIENCE IT IN IMAX

Use Case - Real-Time Health Monitoring for Chronic Disease Management

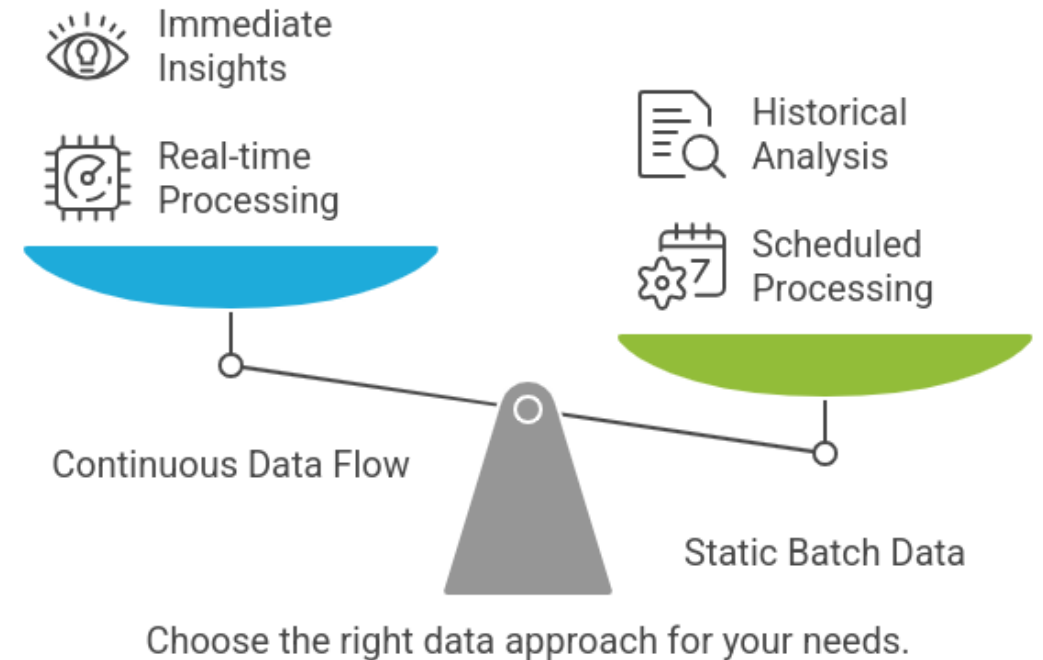
- How stream processing can transform patient care by enabling real-time health monitoring, particularly for chronic diseases such as heart failure or diabetes.



- Source: arxiv.org and confluent.io

Batch Vs Stream

Aspect	Stream Processing	Batch Processing
Latency	Milliseconds to seconds	Minutes to hours
Data Nature	Unbounded, continuous	Bounded, discrete
Processing	Incremental (micro-batches) or event-by-event	Processed all at once
Use Cases	Real-time monitoring, IoT, fraud detection	Reporting, analytics, historical analysis





Apache Spark Structured Streaming

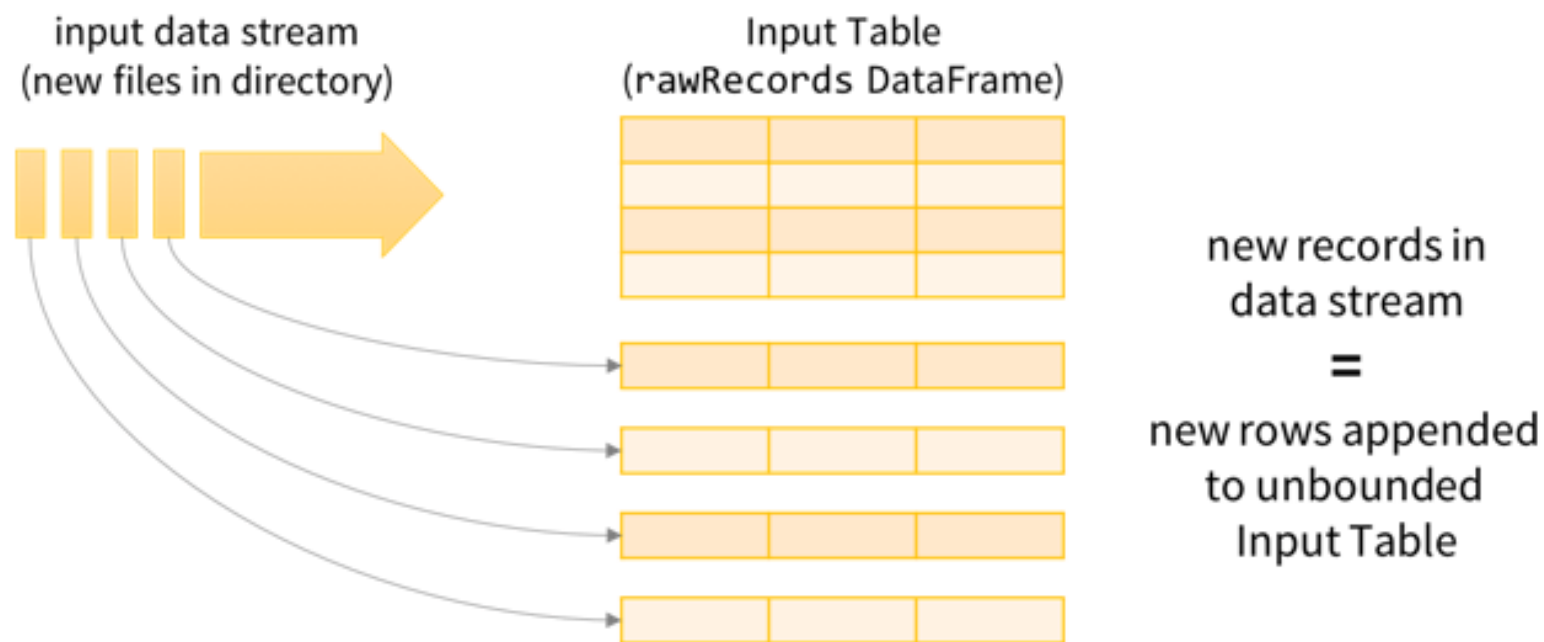
Apache Spark Streaming



Structured Streaming




Every data item that is arriving on the stream is like a new row being appended to the input table.



Structured Streaming Model

Read Stream



Processing pipeline



Write Stream



The possible input sources

- File source
 - Reads files written in a directory as a stream of data.

File-0.json

File-1.json

File-2.json

File-3.json

File-4.json

File-n.json

Files will be processed in the order of file modification time.

If *latestFirst* is set, order will be reversed

Supported file formats are text, CSV, JSON, ORC, Parquet.

Input Stream (file source)

```
streamingInputDF = (  
    spark  
        .readStream                                # it can be “read” for defining static  
processing  
        .schema(jsonSchema)                        # Set the schema of the JSON  
data  
        .option("maxFilesPerTrigger", 1) # Treat a sequence of files as a  
stream by picking one file at a time  
        .json(inputPath)  
)
```

Processing pipeline (examples)

You can do a wide variety of processing/transformation over the incoming data stream.

Split the lines into words

```
words = streamingInputDF.select(  
    explode(  
        split(streamingInputDF.value, " ")  
    ).alias("word") # name of the new field  
)
```

Generate running word count

```
wordCounts = words.groupBy("word").count()
```


Processing pipeline (examples)

```
streamingCountsDF = (  
    streamingInputDF  
        .groupBy(  
            streamingInputDF.action,  
            window(streamingInputDF.time, "1 hour"))  
        .count()  
    )
```

Processing pipeline (examples)

```
result_df = streamingInputDF \  
    .groupBy("name") \  
    .agg({"price": "mean"})
```

The possible output modes

1

Complete Mode: The entire updated result table is written to external storage.

2

Append Mode: Only new rows appended in the result table since the last trigger are written to external storage.

3

Update Mode: Only the rows that were updated in the result table since the last trigger are written to external storage.

Output Sinks

File sink

```
writeStream  
  .format("parquet")      # can be "orc", "json", "csv", etc.  
  .option("path", "path/to/destination/dir")  
  .start()
```

Kafka sink

```
writeStream  
  .format("kafka")  
  .option("kafka.bootstrap.servers", "host1:port1,host2:port2")  
  .option("topic", "updates")  
  .start()
```

Output Sinks ...

Console sink (for debugging) - Prints the output to the console/stdout every time there is a trigger. Both, Append and Complete output modes, are supported.

```
writeStream  
  .format("console")  
  .start()
```

Memory sink (for debugging) - The output is stored in memory as an in-memory table. Both, Append and Complete output modes, are supported.

```
writeStream  
  .format("memory")  
  .queryName("tableName")  
  .start()
```

Output Stream Example

Start running the query that prints the running counts to the console

```
query = wordCounts \  
    .writeStream \  
    .outputMode("complete") \  
    .format("console") \  
    .start()
```

`query.awaitTermination()` # keeps waiting for the termination of the query.

Output Stream Example ...

```
query = result_df.writeStream \  
    .format("console") \  
    .outputMode("complete") \  
    .start()  
query.awaitTermination()
```

Output Stream Example ...

```
query = (  
    streamingCountsDF  
        .writeStream  
        .format("memory")      # memory = store in-memory table  
        .queryName("counts")   # counts = name of the in-memory table  
        .outputMode("complete") # complete = all the counts should be in  
the table  
        .start()  
)
```

References

- <https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>
- <https://docs.databricks.com/structured-streaming/examples.html>
- <https://www.databricks.com/spark/getting-started-with-apache-spark/streaming>
- <https://docs.databricks.com/getting-started/streaming.html#notebook-stream>
- https://www.splunk.com/en_us/blog/learn/stream-processing.html
- <https://arxiv.org/pdf/1707.04364>
- <https://www.confluent.io/en-gb/blog/single-patient-view/>

Workshop

- Defining an appropriate schema for parsing incoming messages.
- Implementing input data stream.
- Designing a pipeline to process data.
- Creating a right stream writer